

101% Job-Ready

DATA ANALYTICS COURSE

Complete Reference Guide

Phase 1 Excel + Statistics + AI	Phase 2 SQL + BigQuery + BI Patterns	Phase 3 Python + ML + Streamlit
Phase 4 Tableau + Power BI + Looker	Phase 5 Cloud + dbt + LLMs	Phase 6 Career Launch + Interviews

16 Weeks

Full curriculum

6 Phases

Theory + code

20+ Parts

Quizzes + practice

5 Projects

Portfolio-ready

Table of Contents

Phase 1 — Excel, Statistics and AI Tools

- P1.1 — Excel Fundamentals + Microsoft Copilot
- P1.2 — Statistics for Analysts
- P1.3 — Advanced Excel + AI Integration
- P1.4 — Phase 1 Capstone

Phase 2 — SQL, BigQuery and Business Analytics

- P2.1 — SQL Fundamentals
- P2.2 — Advanced SQL + Window Functions
- P2.3 — BigQuery + Gemini AI + RFM/Cohort/Funnel
- P2.4 — E-Commerce Analytics Capstone

Phase 3 — Python for Analytics + Machine Learning

- P3.1 — Python Fundamentals + AI Pair Programming
- P3.2 — Exploratory Data Analysis + Visualisation
- P3.3 — Data Wrangling + Scikit-learn + SHAP
- P3.4 — Customer Intelligence Platform Capstone

Phase 4 — BI Tools — Tableau, Power BI, Looker

- P4.1 — Tableau + Einstein AI / Tableau Pulse
- P4.2 — Power BI + Microsoft Copilot + DAX
- P4.3 — ThoughtSpot + Looker + Executive Capstone

Phase 5 — Cloud, Data Engineering and AI Ops

- P5.1 — GCP + BigQuery + dbt + Apache Airflow
- P5.2 — LLMs + RAG + LangChain + Evidently + Capstone

Phase 6 — Career Launch

- P6.1 — Portfolio + CV + LinkedIn + Interviews + Salary

Excel, Statistics and AI Tools

Weeks 1-3 · Parts 1-4

Key Concepts

Excel Fundamentals

- Pivot Tables: drag fields to Rows, Columns, Values, Filters — summarise thousands of rows instantly
- XLOOKUP(lookup_value, lookup_array, return_array) — replaces VLOOKUP, works left and right
- SUMIFS / COUNTIFS / AVERAGEIFS — aggregate with multiple conditions
- Power Query: connect, clean, and reshape data before it enters Excel — every step is recorded
- Named Ranges and structured table references: =SalesTable[Revenue] instead of =B2:B1000

Microsoft Copilot in Excel

- Highlight data > Copilot > Analyse — generates instant insights and chart suggestions
- Ask: 'What are the top 3 trends in this data?' — Copilot summarises in plain English
- Generate formulas: 'Write a formula to calculate year-over-year growth for each row'
- Copilot can add columns with complex formulas, spot anomalies, and create pivot tables on request

Statistics for Analysts

Metric	Definition	When to Use
Mean	Average of all values	Skewed by outliers — use median for skewed data
Median	Middle value when sorted	Best for skewed distributions like revenue
Mode	Most frequent value	Useful for categorical data (most popular product)
Std Dev	Spread around the mean	Low = consistent; High = volatile data
Percentile	Value below which X% falls	Median = 50th percentile
Correlation r	Linear relationship strength (-1 to +1)	Above 0.7 = strong; above 0.9 = very strong

AI SHORTCUT: Paste your data summary into ChatGPT: 'Here is my dataset summary: [paste]. What statistical tests should I run? What are the key distributions? Are there any red flags?'

SQL, BigQuery and Business Analytics

Weeks 4-6 · Parts 1-4

SQL Fundamentals

Core Clauses and Execution Order

Written order: SELECT → FROM → JOIN → WHERE → GROUP BY → HAVING → ORDER BY → LIMIT

Execution order: FROM → JOIN → WHERE → GROUP BY → HAVING → SELECT → ORDER BY → LIMIT

```
-- Template: complete SELECT with all clauses
SELECT region, SUM(revenue) AS total_revenue,
COUNT(*) AS order_count
FROM orders
JOIN customers USING (customer_id)
WHERE status = 'Completed'
AND order_date >= '2024-01-01'
GROUP BY region
HAVING SUM(revenue) > 10000
ORDER BY total_revenue DESC
LIMIT 10;
```

JOIN Types

Join Type	Returns	Use When
INNER JOIN	Rows matching in BOTH tables	You only want clean matched records
LEFT JOIN	All from left + matches from right (NaN where no match)	Keep all transactions, add dimension data
RIGHT JOIN	All from right + matches from left	Rarely used — rewrite as LEFT with tables swapped
FULL OUTER	All rows from BOTH tables	Full audit — find unmatched records on either side

Window Functions

```

-- ROW_NUMBER: latest order per customer
ROW_NUMBER() OVER (PARTITION BY customer_id ORDER BY order_date DESC)

-- RANK vs DENSE_RANK
RANK() OVER (ORDER BY revenue DESC) -- skips numbers after ties: 1,1,3
DENSE_RANK() OVER (ORDER BY revenue DESC) -- no skips: 1,1,2

-- LAG: month-over-month comparison
LAG(revenue, 1) OVER (ORDER BY month) -- previous month's revenue

-- Running total
SUM(revenue) OVER (ORDER BY order_date ROWS UNBOUNDED PRECEDING)

-- % of total per group
revenue / SUM(revenue) OVER (PARTITION BY region) * 100

```

Business Analytics Patterns

RFM Segmentation

```

WITH rfm AS (
  SELECT customer_id,
    DATEDIFF(CURRENT_DATE, MAX(order_date)) AS recency_days,
    COUNT(*) AS frequency,
    SUM(revenue) AS monetary
  FROM orders WHERE status='Completed'
  GROUP BY customer_id
)
SELECT *,
  NTILE(4) OVER (ORDER BY recency_days ASC) AS r_score,
  NTILE(4) OVER (ORDER BY frequency DESC) AS f_score,
  NTILE(4) OVER (ORDER BY monetary DESC) AS m_score
FROM rfm;

-- Segments: r=4 AND f>=3 = Champion
-- r<=2 AND f>=3 = At Risk
-- r=1 AND f=1 = Lost

```

Cohort Retention

```

-- Step 1: cohort month = customer's first order month
DATE_TRUNC(MIN(order_date), MONTH) AS cohort_month

-- Step 2: months since first order
DATE_DIFF(order_month, cohort_month, MONTH) AS month_number

-- Step 3: retention rate
retained_customers * 100.0 / cohort_size AS retention_pct

-- Benchmarks: e-commerce Month 1 = 20-40%, Month 3 = 15-25%

```

BigQuery Key Differences

Feature	BigQuery Syntax	Standard SQL
Table reference	`project.dataset.table`	table_name
Date truncation	DATE_TRUNC(date, MONTH)	DATETRUNC(MONTH, date)

String agg	STRING_AGG(col, ',')	GROUP_CONCAT(col)
Current date	CURRENT_DATE()	GETDATE() or NOW()

Python for Analytics and Machine Learning

Weeks 7-10 · Parts 1-4

Pandas Essentials

```
import pandas as pd
import numpy as np

df = pd.read_csv('orders.csv')

# First look (run these on every new dataset)
df.shape # (rows, cols)
df.info() # types + null counts
df.describe() # stats – check 50% (median) vs mean for skew
df.isnull().sum() # null counts per column

# Filtering (boolean indexing)
completed = df[df['status'] == 'Completed']
north_big = df[(df['region']=='North') & (df['revenue']>500)]

# GroupBy aggregation
df.groupby('region').agg(
    total_revenue = ('revenue', 'sum'),
    order_count = ('order_id', 'count'),
    avg_order = ('revenue', 'mean'),
).reset_index()

# Derived columns
df['revenue'] = df['units'] * df['price']
df['margin_pct'] = (df['revenue'] - df['cost']) / df['revenue'] * 100
df['is_returned'] = (df['status'] == 'Returned').astype(int)
df['is_weekend'] = df['order_date'].dt.dayofweek.isin([5,6]).astype(int)
```

7-Step EDA Framework

1. Shape and Schema — `df.shape`, `df.info()`, column types, row count vs source
2. Missing Values — `df.isnull().sum() / len(df)`. More than 20% is a red flag
3. Distributions — histograms, skewness. Right-skewed? Use median not mean
4. Descriptive Statistics — `df.describe()`. Compare mean vs 50% for skew
5. Categorical Summaries — `value_counts()`, `unique()` before any groupby
6. Correlations — `df.corr()`, `sns.heatmap()`. $r > 0.9$ = multicollinearity
7. Outlier Detection — IQR rule: below $Q1 - 1.5 * IQR$ or above $Q3 + 1.5 * IQR$

Scikit-learn ML Pipeline

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, roc_auc_score

X = df[['recency_days', 'frequency', 'monetary', 'avg_order_val']]
y = df['churned']

X_train, X_test, y_train, y_test = train_test_split(
X, y, test_size=0.2, random_state=42, stratify=y)

scaler = StandardScaler()
X_tr = scaler.fit_transform(X_train) # fit on train ONLY
X_te = scaler.transform(X_test) # transform test with same params
rf = RandomForestClassifier(n_estimators=200, random_state=42)
rf.fit(X_tr, y_train)
y_pred = rf.predict(X_te)
print(roc_auc_score(y_test, y_pred)) # target: above 0.80
print(classification_report(y_test, y_pred))

```

SHAP Explainability

```

import shap
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(X_te)

# Global: which features drive predictions most?
shap.summary_plot(shap_values[1], X_te, plot_type='bar')

# Individual: why did this customer get a high churn score?
# Positive SHAP = pushes prediction UP (toward churn)
# Negative SHAP = pushes prediction DOWN (away from churn)
shap.force_plot(explainer.expected_value[1],
shap_values[1][0], X_te[0])

```


PHASE 4

BI Tools — Tableau, Power BI, Looker, ThoughtSpot

Weeks 11-13 · Parts 1-3

Tableau

Calculated Fields and LOD Expressions

```
-- Basic calculated field
Margin % = ([Revenue] - [Cost]) / [Revenue]

-- Conditional (like Excel IF)
IF [Revenue] >= 1000 THEN 'High'
ELSEIF [Revenue] >= 500 THEN 'Medium'
ELSE 'Low' END

-- LOD: FIXED — compute at region level, ignore view filters
{ FIXED [Region] : SUM([Revenue]) }

-- LOD: % of total
SUM([Revenue]) / { FIXED : SUM([Revenue]) }

-- LOD: customer's first order date (for cohort analysis)
{ FIXED [Customer ID] : MIN([Order Date]) }

-- Parameter: dynamic metric switcher
CASE [Select Metric]
WHEN 'Revenue' THEN SUM([Revenue])
WHEN 'Profit' THEN SUM([Profit])
WHEN 'Orders' THEN COUNTD([Order ID])
END
```

Power BI — DAX

```

-- Core measures
Total Revenue = SUM(orders[revenue])
Return Rate = DIVIDE(CALCULATE(COUNTROWS(orders),
orders[status]='Returned'), COUNTROWS(orders), 0)

-- CALCULATE: the most important DAX function
North Revenue = CALCULATE([Total Revenue], orders[region]='North')
% of Total = DIVIDE([Total Revenue],
CALCULATE([Total Revenue], ALL(orders)), 0)

-- Time intelligence (requires a Date table)
Revenue YTD = CALCULATE([Total Revenue],
DATESYTD(DateTable[Date]))
Revenue MoM % = VAR curr = [Total Revenue]
VAR prev = CALCULATE([Total Revenue],
DATEADD(DateTable[Date],-1,MONTH))
RETURN DIVIDE(curr - prev, prev, 0)

-- RANKX: rank products by revenue
Product Rank = RANKX(ALL(orders[product]),
[Total Revenue], , DESC, Dense)

```

BI Tool Comparison

Tool	Cost	Strength	Best For
Tableau	\$70/mo	Complex visuals, LOD, custom analysis	Complex interactive dashboards
Power BI	Free+\$10	Microsoft integration, DAX, data model	Microsoft 365 organisations
Looker Studio	Free	Google ecosystem, shareable links	Quick shareable dashboards
ThoughtSpot	Enterprise	Search-driven, Spotter AI	Self-service analytics at scale

Dashboard Design Principles

- One question per dashboard — everything on the page answers that question only
- KPIs first (top-left), context second (charts below), details last (tables at bottom)
- Filter actions over separate dashboards — click a region to filter all charts at once
- Remove everything that does not earn its place — gridlines, borders, unnecessary legends
- Consistent colour with meaning — red = bad, green = good, grey = neutral/context
- Design for your audience — executive wants 3 numbers, analyst wants 6 charts

Cloud, Data Engineering and AI Ops

Weeks 14-15 · Parts 1-2

The Modern Data Stack

Layer	Tools	Purpose
01 Sources	CRM, App DB, APIs, Spreadsheets	Where data is born
02 Ingestion	Fivetran, Airbyte, Stitch	Move raw data into warehouse
03 Warehouse	BigQuery, Snowflake, Redshift	Central analytics store
04 Transform	dbt (SQL models)	Clean, model, enrich data
05 Orchestrate	Apache Airflow, Prefect, Dagster	Schedule and monitor pipelines
06 Semantic	dbt metrics, Looker LookML	Define business metrics once
07 BI	Tableau, Power BI, Looker Studio	Visualise for stakeholders

dbt Core Concepts

```
-- Staging model: clean raw data (materialized='view')
{{ config(materialized='view') }}
SELECT order_id,
UPPER(TRIM(region)) AS region,
order_date::DATE AS order_date,
units * price AS revenue
FROM {{ source('raw', 'orders') }}
WHERE order_id IS NOT NULL;

-- Mart model: business-ready table (materialized='table')
{{ config(materialized='table') }}
SELECT o.*, c.country, p.category
FROM {{ ref('stg_orders') }} o
LEFT JOIN {{ ref('stg_customers') }} c USING (customer_id)
LEFT JOIN {{ ref('stg_products') }} p USING (product_id);

-- Data quality tests in schema.yml
-- unique, not_null, accepted_values, relationships

-- Key commands:
-- dbt run -- build all models
-- dbt test -- run quality tests
-- dbt docs serve -- open documentation in browser
```

Apache Airflow — DAG Pattern

```

from airflow import DAG
from airflow.operators.bash import BashOperator
from datetime import datetime, timedelta

with DAG('retailco_daily', schedule_interval='0 6 * * *',
start_date=datetime(2024,1,1), catchup=False,
default_args={'retries': 2,
'retry_delay': timedelta(minutes=5)}) as dag:
upload = BashOperator(task_id='upload_to_gcs', ...)
load = BashOperator(task_id='load_to_bq', ...)
run = BashOperator(task_id='dbt_run', bash_command='dbt run')
test = BashOperator(task_id='dbt_test', bash_command='dbt test')

upload >> load >> run >> test # execution order

```

LLMs and RAG for Analytics

```

# RAG pipeline: LangChain + FAISS + OpenAI
from langchain_openai import ChatOpenAI, OpenAIEmbeddings
from langchain_community.vectorstores import FAISS
from langchain.chains import RetrievalQA
from langchain.schema import Document

# 1. Convert data rows to Documents
docs = [Document(page_content=row.to_string())
for _, row in df.iterrows()]

# 2. Embed into vector store
vectordb = FAISS.from_documents(docs, OpenAIEmbeddings())

# 3. Build Q&A chain
qa = RetrievalQA.from_chain_type(
    llm=ChatOpenAI(model='gpt-4o-mini', temperature=0),
    retriever=vectordb.as_retriever(search_kwargs={'k': 5}))

# 4. Ask questions
print(qa.invoke('Which customers are at highest churn risk?'))

```

KEY INSIGHT — RAG vs Plain LLM: Plain LLMs answer from training data and may hallucinate. RAG retrieves relevant chunks from YOUR data first, then generates grounded answers. Always use `temperature=0` for analytics tasks — deterministic, consistent responses.

PHASE 6

Career Launch — Get Hired as a Data Analyst

Week 16

Portfolio Projects Summary

Phase	Project Title	Tools	Live Link
Phase 1	Excel Dashboard	Excel, Copilot, Power Query	Tableau Public or Google Drive
Phase 2	SQL Analytics Engine	BigQuery, SQL, Looker Studio	GitHub + Looker Studio URL
Phase 3	Customer Intel Platform	Python, Scikit-learn, Streamlit	GitHub + Streamlit Cloud URL
Phase 4	Executive Analytics Suite	Tableau, Power BI, Looker	Tableau Public + Power BI Service
Phase 5	AI Analytics Platform	dbt, Airflow, LangChain	GitHub + Streamlit chatbot URL

CV Writing — The Magic Formula

[Action Verb] + [what you did] + [using what tool/method] + [measurable outcome] Example: 'Automated monthly sales reporting using Python and Apache Airflow, reducing analyst time from 8 hours to 15 minutes per month.'

- Start every bullet with a strong verb: Built, Automated, Designed, Reduced, Increased, Delivered
- Every bullet must have a number — %, £/\$, hours, rows, customers, % reduction
- ATS tip: mirror exact keywords from the job description. If JD says 'Tableau', use 'Tableau'
- Projects section is critical for career changers — this is your evidence of skill
- 1 page for under 5 years experience. 2 pages maximum for senior roles. Always PDF.

LinkedIn Headline Formula

WEAK: 'Student at University of Manchester' STRONG: 'Data Analyst | SQL · Python · Tableau · Power BI | Turning data into decisions'

Interview — Top 10 Questions You MUST Prepare

[SQL] Write a query to find the top 3 customers by revenue in each region

ROW_NUMBER() OVER (PARTITION BY region ORDER BY revenue DESC) — filter WHERE rank <= 3

[SQL] What is the difference between WHERE and HAVING?

WHERE filters rows before aggregation. HAVING filters groups after GROUP BY. Use HAVING when filtering on an aggregate like SUM(revenue) > 1000.

[Python] How would you handle 30% null values in a key column?

Investigate WHY they are null first. If random: fill with median/mode. If systematic: investigate the pipeline. Never auto-drop.

[BI] What is CALCULATE() in DAX?

CALCULATE(expression, filter1...) evaluates a measure with modified filter context. Most important DAX function — powers % of total, context-aware measures.

[ML] What is overfitting? How do you detect it?

Overfitting: model memorises training data, fails on new data. Detect with: train vs test accuracy gap, cross-validation. Prevent: regularisation, simpler model.

[Behavioural] Tell me about a time you found an insight that changed a decision.

Use STAR: Situation, Task, Action (specific steps you took), Result (quantified outcome). Prepare 4-5 specific stories from your portfolio projects.

[Case Study] Revenue dropped 15% last month. Walk me through your investigation.

(1) Verify drop is real — check for data pipeline issues. (2) Segment: by region, product, channel. (3) Find the 80% driver. (4) Hypothesise cause. (5) Quantify.

[LOD] What is a LOD expression in Tableau?

Level of Detail expression — computes a metric at a different granularity than the view. FIXED ignores view filters. Example: {FIXED [Region]: SUM([Revenue])}

[dbt] What does ref() do in dbt?

ref('model_name') creates a dependency — dbt builds stg_orders before any model that references it. Also resolves full table name per environment automatically.

[RAG] What is the difference between a vector database and a traditional database?

Traditional DB: exact match search. Vector DB: semantic similarity search using embeddings. Powers RAG pipelines — finds relevant documents by meaning, not keywords.

Salary — 2025 Market Rates

Level	UK	US	Global Remote
Junior (0-2 yrs)	£28-40K	\$60-85K	\$50-70K
Mid (2-5 yrs)	£40-60K	\$85-120K	\$70-100K
Senior (5+ yrs)	£60-90K	\$120-180K	\$100-150K
Analytics Engineer	£55-85K	\$130-190K	\$110-160K

NEGOTIATION RULE: Never accept on the spot. Always say: 'Thank you — I am genuinely excited. Can I have until [date 3-5 days away] to review the full package?' The first person to name a number loses leverage. Make them go first.

Quick Reference — Full Tool Stack

SQL

- SELECT / FROM / WHERE / GROUP BY / HAVING / ORDER BY
- INNER / LEFT / RIGHT / FULL OUTER JOIN
- Window functions: ROW_NUMBER, RANK, DENSE_RANK, LAG, LEAD, SUM OVER
- CTEs: WITH cte AS (SELECT ...) SELECT * FROM cte
- Subqueries: scalar, IN, derived table, EXISTS
- BigQuery: `project.dataset.table` backtick references

Python / Pandas

- pd.read_csv / to_csv / read_excel / to_excel
- df[df['col'] == 'value'] — boolean indexing
- df.groupby('col').agg(metric=('col','sum')).reset_index()
- pd.merge(left, right, on='key', how='left')
- pd.pivot_table(df, values, index, columns, aggfunc)
- pd.melt(df, id_vars, value_vars, var_name, value_name)
- np.where(condition, if_true, if_false)

Tableau

- Dimensions (blue) = categorical. Measures (green) = numeric
- LOD: { FIXED [dim] : AGG([measure]) }
- Parameters: user-controlled variable for dynamic calculations
- Filter action: Dashboard > Actions > Filter
- Tableau Pulse: AI auto-narratives and anomaly detection

Power BI / DAX

- Measure: CALCULATE, SUM, COUNTROWS, DIVIDE
- CALCULATE(expr, filter) — override filter context
- ALL(table) — remove all filters from table
- DATEADD / DATESYTD / SAMEPERIODLASTYEAR — time intelligence
- Power Query: M language, visual ETL, Applied Steps

dbt

- {{ ref('model_name') }} — dependency management
- {{ source('schema', 'table') }} — raw table reference
- Materialise: view (no storage), table (stored result)
- Tests: unique, not_null, accepted_values, relationships
- dbt run / dbt test / dbt docs generate + serve

LLMs / RAG

- temperature=0 for deterministic analytics responses
- Embeddings = vector representation of semantic meaning
- RAG = vector search + LLM generation grounded in YOUR data
- LangChain: RetrievalQA, FAISS, PromptTemplate
- Evidently: drift monitoring — retrain if >30% features drift

AI Prompt Library — Best Prompts from the Course

SQL Generation

RFM Query

I use BigQuery. My table is orders(customer_id, order_date, units, price, status). Write a complete RFM segmentation query using NTILE(4) for R, F, and M scoring. Label customers as Champion, Loyal, At Risk, or Lost using CASE WHEN. Comment every step.

Cohort Analysis

Write a BigQuery cohort retention analysis. Table: orders(customer_id, order_date, status). Use DATE_TRUNC for monthly cohorts, DATE_DIFF for month numbers, and calculate retention_pct. Save as a view vw_cohort_retention.

Funnel Analysis

I have a user_events table: user_id, event_type (page_view/product_view/add_to_cart/checkout/purchase), event_date. Write a funnel query for last 30 days showing users at each stage, % of top, and step-by-step drop-off %.

Python / Pandas

Data Cleaning Pipeline

I have a CSV with columns: order_id, customer_id, order_date (string), units, price, cost, status, region. Write pandas code to: cast order_date to datetime, strip whitespace from region, calculate revenue = units * price and profit = revenue - cost, remove duplicates on order_id, report null counts.

Churn Model

Build a complete Python churn prediction pipeline: load customer_features.csv (recency_days, frequency, monetary, avg_order_value, churned), split 80/20, scale with StandardScaler, train Random Forest, print ROC-AUC and classification_report, and list feature importances.

EDA Report

I completed EDA on a retail dataset. Findings: [paste stats]. Write a professional EDA report: Executive Summary, Data Quality, Distributions, Correlations, Anomalies, Next Steps. Include specific numbers in every bullet.

DAX / Power BI

Time Intelligence Pack

Write DAX measures for Power BI. Table 'orders' columns: revenue, status, order_date. I need:
(1) Total Revenue for Completed orders, (2) Revenue MoM % change, (3) Revenue YTD, (4) Return Rate %, (5) Revenue % of grand total ignoring all filters.

LangChain / RAG

Analytics Chatbot

Write a LangChain RetrievalQA pipeline that: (1) loads customer_features.csv, (2) converts each row to a LangChain Document, (3) builds a FAISS vector store with OpenAIEmbeddings, (4) creates a RetrievalQA with a custom prompt that says 'Answer using ONLY the context provided', (5) answers 5 sample business questions.

Career

CV Bullet Points

Help me write 4 CV bullet points for this project: [describe your project, tools, problems solved, results]. Each bullet: starts with a strong action verb, includes the specific tool, ends with a measurable outcome. Use STAR format.

Interview Mock

Act as a senior data analyst interviewing me for a data analyst role. Ask me one technical question about [topic]. After my answer: (1) score out of 10, (2) what I did well, (3) what to improve, (4) a stronger version of my answer.



You Are Now
101% Job-Ready.
SQL · Python · Tableau · Power BI · Looker · BigQuery
dbt · Airflow · LangChain · SHAP · Streamlit · OpenAI API

101% Job-Ready Data Analytics Course · 16 Weeks · 6 Phases · 5 Portfolio Projects